

1.4: Analysis of Algorithms

Exercise 1. Show that the number of different triples that can be chosen from n items is precisely $n(n-1)(n-2)/6$. *Hint:* Use mathematical induction or a counting argument.

Solution.

Proof. This is the problem of choosing a combination of 3 out of n , which is given by $\binom{n}{3}$, and

$$\binom{n}{3} = \frac{n!}{3!(n-3)!} = \frac{n \cdot (n-1) \cdot (n-2) \cdot (n-3)!}{3!(n-3)!} = \frac{n(n-1)(n-2)}{6}$$

□

Exercise 2. Modify `ThreeSum` to work properly even when the `int` values are so large that adding two of them might cause integer overflow.

Solution. There are two cases when it comes to overflow:

- (i) *Positive overflow.* The sum exceeds `Integer.MAX_VALUE`. If two terms sum to `Integer.MAX_VALUE + 1`, overflow occurs, and the value wraps around to `Integer.MIN_VALUE`. Thus, if $a + b = \text{Integer.MAX_VALUE} + 1$ and $c = \text{Integer.MIN_VALUE}$, then we have a valid sum. However, if $a + b$ sums to anything larger, then no value of c will do because c cannot be smaller than `Integer.MIN_VALUE`.
- (ii) *Negative overflow.* The sum of two negative numbers a and b , yielding a value below `Integer.MIN_VALUE`. In this case, it's impossible to have $a + b = c$ for any 32-bit two's complement integer c .

See the `com.segarciat.algs4.ch1.sec4.ex02.ThreeSum` class.

Exercise 3. Modify `DoublingTest` to use `StdDraw` to produce plots like the standard and log-log plots in the text, rescaling as necessary so that the plot always fills a substantial portion of the window.

Solution. See my `com.segarciat.algs4.ch1.sec4.ex03.DoublingTest` class.

Exercise 4. Develop a table like the one on page 181 for `TwoSum`.

Solution. The `TwoSum` program referenced is:

```
public class TwoSum {
    public static int count(int[] a)
    {    // Count pairs that sum to 0;
        // A: The entire method body
        int n = a.length;
```

```

    int count = 0;
    // B: The outermost for loop and its body (not including the statement i
        = 0)
    for (int i = 0; i < n; i++)
        // C: The innermost for loop and its body (not including the
            statement j = i + 1)
        for (int j = i + 1; j < n; j++)
            if (a[i] + a[j] == 0)
                // D: The number of times the if block is executed
                count++;
    return count;
}
}

```

The table in 181 is used to analyze the running time of by keeping track of the frequency of each statement block, that is, the number of times the block is executed:

Statement block	Time in seconds	Frequency	Total time
D	t_0	x (depends on input)	t_0x
C	t_1	$\binom{n}{2}$	$t_1 \binom{n}{2}$
B	t_2	n	t_2n
A	t_3	1	t_3

Therefore, the grand total is:

$$\begin{aligned}
 t_1 \binom{n}{2} + t_2n + t_3 + t_0x &= t_1 \frac{n(n-1)}{2} + t_2n + t_3 + t_0x \\
 &= \frac{t_1}{2}n^2 + \left(t_2 - \frac{t_1}{2}\right)n + t_3 + t_0x
 \end{aligned}$$

The tilde approximation (assuming x is small) is

$$\sim \left(\frac{t_1}{2}\right)n^2$$

Hence the order of growth is n^2 .

Exercise 5. Give tilde approximations for the following quantities:

- (a) $n + 1$
- (b) $1 + 1/n$
- (c) $(1 + 1/n)(1 + 2/n)$
- (d) $(2n^3 - 15n^2)$
- (e) $\lg(2n)/\lg n$
- (f) $\lg(n^2 + 1)/\lg n$
- (g) $n^{100}/2^n$

Solution. The definition of tilde approximation given on page 179 of [SW11] says $\sim f(n)$ represents a function that, when divided by $f(n)$, approaches 1 as n grows.

- (a) $\sim n$
- (b) ~ 1 , since $1/n$ approaches 0 as n grows, and hence $1 + 1/n$ approaches 1 as n grows.
- (c) ~ 1 , similar to (b).
- (d) $\sim 2n^3$
- (e) Since $\lg(2n) = \lg(2) + \lg(n)$, this means $\lg(2n)/\lg n = \lg(2)/\lg(n) + 1$. Hence, this is again ~ 1 .
- (f) Since $\lg(n^2 + 1) \approx \lg(n^2)$, and $\lg(n^2)/\lg(n) = 2\lg(n)/\lg(n)$, we conclude that this is ~ 2 .
- (g) Here we can just say $\sim (n^{100}2^{-n})$.

Exercise 6. Give the order of growth (as a function of n) of the running times of each of the following code fragments:

(a)

```
int sum = 0;
for (int k = n; k > 0; k /= 2)
    for (int i = 0; i < k; i++)
        sum++;
```

(b)

```
int sum = 0;
for (int i = 1; i < n; i *= 2)
    for (int j = 0; j < i; j++)
        sum++;
```

(c)

```
int sum = 0;
for (int i = 1; i < n; i *= 2)
    for (int j = 0; j < n; j++)
        sum++;
```

Solution. Let $m = \lfloor \lg n \rfloor + 1$. Then m is the number of bits needed to represent n in binary, and is the frequency of execution of the outer loop. In particular, $2^{m-1} \leq n < 2^m$.

- (a) For each value of k , the i loop block containing the statement `sum++` will execute k times, where $k = \lfloor n/2^j \rfloor$ for $0 \leq j \leq m$. Thus the total number of times that

`sum++` is executed is approximately given by

$$\begin{aligned}
\sum_{j=0}^{m-1} \left\lfloor \frac{n}{2^j} \right\rfloor &< \sum_{j=0}^{m-1} \left\lfloor \frac{2^m}{2^j} \right\rfloor \\
&= \sum_{j=0}^{m-1} \frac{2^m}{2^j} \\
&= 2^m \sum_{j=0}^{m-1} \frac{1}{2^j} \\
&= 2^m \cdot \left(2 - \frac{1}{2^{m-1}} \right) \\
&= 2^{m+1} - 2 \\
&= 2(2^m - 1) \\
&\leq 2(2^{\lg n + 1} - 1) \\
&= 2 \cdot (2n - 1)
\end{aligned}$$

The order of growth is n , or linear.

- (b) This is similar to (a), but the analysis is much simpler: `i` takes on the values 2^k for $0 \leq k < m$, and the `j` loop executes `i` times for each `i`. Thus the total number of times `sum++` runs is approximately:

$$\begin{aligned}
\sum_{k=0}^{m-1} 2^k &= 2^m - 1 = 2^{\lg n + 1} - 1 \\
&\leq 2^{\lg n + 1} - 1 \\
&= 2n - 1
\end{aligned}$$

Thus the order of growth is n , or linear.

- (c) The outer loop executes m times, so again `i` takes on the values 2^k for $0 \leq k < m$. For each `i`, the `j` loops executes `n` times. Thus the total number of times that `sum++` executes is nm times, which is exactly $n(\lfloor \lg n + 1 \rfloor)$ times. Hence the order of growth is linearithmic.

Exercise 7. Analyze `ThreeSum` under a cost model that counts arithmetic operations (and comparisons) involving the input numbers.

Solution. The arithmetic operations in `ThreeSum` are all additions. The comparison operations are less than (`<`) and equal to (`==`).

Altogether, there are $n + 1$ comparisons from `i < n`, $\binom{n}{2} + 1$ comparisons from `j < n` (once for each pair (i, j)), $\binom{n}{3} + 1$ comparisons from `k < n`, and $\binom{n}{3}$ comparisons for `a[i] + a[j] + a[k] == 0` (one for each triple (i, j, k)).

There are n additions from `i++`, $\binom{n}{2}$ additions for `j++`, $\binom{n}{3}$ additions for `k++`, $2 \cdot \binom{n}{3}$ additions for `a[i] + a[j] + a[k]` (since it involves two additions for each triple), and the number of times `count++` executes is indeterminate since it depends precisely on how many times the control expression of the `if` evaluates to `true`.

Overall, then, the cost is dominated by the statements in the innermost `k` loop. Therefore, we can say that `ThreeSum` uses about $5\binom{n}{3}$ or $\sim \frac{5}{6}n^3$ arithmetic operations (including comparisons), and hence its order of growth is cubic under a cost model that counts arithmetic operations and comparisons.

Exercise 8. Write a program to determine the number of pairs of values in an input file that are equal. If your first try is quadratic, think again and use `Arrays.sort()` to develop a linearithmic solution.

Solution. See the `com.segarciat.algs.ch1.sec4.ex08.EqualNumberPairs` class. It uses `Arrays.sort()`, which has a linearithmic order of growth. Then it uses the fact that equal numbers are adjacent to each other to compute the frequency of occurrence of each number. If a number has a frequency f , then there are $\binom{f}{2}$ equal pairs corresponding to that number. This loop operates in linear time since it performs at most a constant number of operations in each iteration, and it iterates n times, where n is the number of items in the file.

Exercise 9. Give a formula to predict the running time of a program for a problem of size n when doubling experiments have shown that the doubling factor is 2^b and the running time for problems of size n_0 is t .

Solution. Let $T(n)$ be the running time of the program. Since the doubling factor is 2^b , we know that $T(n)$ has an order of growth approximately n^b , as claimed in [SW11] on page 192 (Section 1.4). Hence,

$$\begin{aligned} T(n) &\sim n^b \\ &= \left(\frac{n}{n_0}\right)^b n_0^b \\ &= \left(\frac{n}{n_0}\right)^b n_0^b \\ &\sim \left(\frac{n}{n_0}\right)^b T(n_0) \end{aligned}$$

Since $T(n_0) = t$, we can predict $T(n)$ to be approximately $\left(\frac{n}{n_0}\right)^b t$.

Exercise 10. Modify binary search so that it always returns the element with the smallest index that matches the search element (and still guarantees logarithmic running time).

Solution. See the `com.segarciat.algs4.ch1.sec4.ex10.BinarySearch` class. I implemented an `indexOf()` method that is similar to my `rank()` implementation in Exercise 1.1.29; see `com.segarciat.algs4.ch1.sec1.ex29.BinarySearch`.

In Exercise 1.1.29, the return value is a number that is between 0 and `a.length`, inclusive. There are two cases:

- (i) When the `key` exists in the array, the return value `a[lo]` has value `key`. Moreover, `lo` is the smallest index such that `a[lo]` equals `key`.

- (ii) When `key` is not in the array, then either `lo` is `a.length` (indicating that `key` exceeds every value in the array) or `lo` is between 0 and `a.length - 1`, but `a[lo]` is not equal to `key`.

My implementation for this exercise adapts the code my `rank()` method in 1.1.29 with these considerations.

Exercise 11. Add an instance method `howMany()` to `StaticSETofInts` (page 99) that finds the number of occurrences of a given key in time proportional to $\log n$ in the worst case.

Solution. My implementation again adapts the code from Exercise 1.1.29, namely the `rank()` and `rankGe()` methods. Both adaptations run in $\log n$ in the worst case. The `howMany()` calls each once, and hence it completes in time proportional to $\log n$ as well.

Exercise 12. Write a program that, given two sorted arrays of n `int` values, prints all elements that appear in both arrays, in sorted order. The running time for your program should be proportional to n in the worst case.

Solution. See the `com.segarciat.algs4.ch1.sec4.ex12.PrintTwoSortedArrays` class.

Exercise 13. Using the assumptions developed in the text, give the amount of memory needed to represent an object of each of the following types:

- (a) `Accumulator`
- (b) `Transaction`
- (c) `FixedCapacityStackOfStrings` with capacity `capacity` and `n` entries.
- (d) `Point2D`
- (e) `Interval1D`
- (f) `Interval2D`
- (g) `Double`

Solution. According to page 201 (Section 1.4) in [SW11], each object has an associated overhead of 16 bytes. Also, memory usage is typically padded to be a multiple of 8 bytes on a 64-bit machine. Hereafter I assume a 64-bit machine:

- (a) An `Accumulator` object requires 16 bytes of overhead, 8 bytes for the `sum` instance variable of type `double`, 4 bytes for the `n` instance variable of type `int`, and 4 bytes of padding. The grand total is 32 bytes.
- (b) A `Transaction` object requires 16 bytes of overhead, 8 bytes for the `who` reference variable of type `String`, 8 bytes for the `when` instance variable of type `Date`, and 8 bytes for the `amount` instance variable of type `double`. These add up to 40 bytes. We also account for the cost plus the cost of a `String` and a `Date` object. For a `String`, assuming Java 7 and later, the cost is $56 + 2n$ bytes (see page 202 on [SW11]), where

n is the number of characters in the string. For a `edu.princeton.cs.algs4.Date` object, the cost is 32 bytes (see page 201 on [SW11]). Altogether, this amounts to $40 + 56 + 2n + 32 = 128 + 2n$ bytes, where n is the number of characters in the `who` instance variable.

- (c) A `FixedCapacityStackOfStrings` object requires 16 bytes of overhead, 4 bytes for the `n` instance variable of type `int`, 8 bytes for the `a` instance variable of type `String[]`, and 4 bytes of padding to be 32 bytes. According to page 202 in [SW11], an array of size `capacity` takes up $24 + 8 \cdot \text{capacity}$ bytes, plus the cost of the `n` string entries. If we let m be the value of `capacity`, this means a total of $56 + 8m$ plus the cost of the `n` objects of type `String`, whose lengths (and hence memory requirements) vary.
- (d) A `Point2D` object requires 16 bytes of overhead, and 8 bytes for each of the instance variables `x` and `y` of type `double`. The grand total is 32 bytes.
- (e) An `Interval1D` object requires 16 bytes of overhead, and 8 bytes for each of the instance variables `min` and `max` of type `double`. The grand total is 32 bytes.
- (f) An `Interval2D` object requires 16 bytes of overhead, 8 bytes for each of the instance `x` and `y` of type `Interval1D`, and 32 bytes for the cost of each `Interval1D` object. The total is 96 bytes.
- (g) A `Double` object requires 16 bytes of overhead and 8 bytes for its instance variable of type `double`. Overall this takes up 24 bytes.

Exercise 1.4.14. Develop an algorithm for the 4-*sum* problem.

Solution. See the `com.segarciat.algs4.ch1.sec4.ex14.FourSum` class. I did not consider overflow.

Exercise 1.4.15. *Faster 3-sum.* As a warmup, develop an implementation `TwoSumFaster` that uses a *linear* algorithm to count the pairs that sum to zero after the array is sorted (instead of the binary-search-based linearithmic algorithm). Then apply a similar idea to develop a quadratic algorithm for the 3-sum problem.

Solution. I implemented `TwoSumFaster` by scanning from opposite sides of the array. I accounted for the possibility that duplicates may exist by including an inner loop to count duplicates and applying the multiplication principle of counting. Overall, the `TwoSumFaster` algorithm is still linearithmic because it uses `Arrays.sort()`, which is linearithmic. However, this still satisfies the constraints of the exercises which requires a linear algorithm *after* sorting. I then implemented `ThreeSumFaster` by applying the algorithm in `TwoSumFaster` a total of n times, where n is the array length, once for each i between 0 and $n - 1$. The algorithm in `ThreeSumFaster` is quadratic: the `Arrays.sort()` is linear, while the `i` loop containing the scan from both ends is quadratic.

Exercise 1.4.16. *Closest pair (in one dimension).* Write a program that, given an array `a[]` of n `double` values, finds a *closest pair*: two values whose difference is no greater than the difference of any other pair (in absolute value). The running time of your program should be linearithmic in the worst case.

Solution. See `com.segarciat.algs4.ch1.sec4.ex16.ClosestPair1D`.

Exercise 1.4.17. *Farthest pair (in one dimension).* Write a program that, given an array `a[]` of n `double` values, finds a *farthest pair*: two values whose difference is no smaller than the difference of any other pair (in absolute value). The running time of your program should be linear in the worst case.

Solution. See the `com.segarciat.algs4.ch1.sec4.ex17.FarthestPair1D` class.

Exercise 1.4.18. *Local minimum of an array.* Write a program that, given an array `a[]` of n distinct integers, finds a *strict local minimum*: an entry `a[i]` that is strictly less than its neighbors. Each internal entry (other than `a[0]` and `a[n-1]`) has 2 neighbors. Your program should use $\sim \lg n$ compares in the worst case.

Solution. First we can prove that there must be a local minimum in an array of distinct values.

Proof. The proof is by contradiction. That is, suppose that no local minimum exists. Then `a[0]` is not a local minimum, so it is larger than its neighbor `a[1]`. Since `a[1]` is not a local minimum, it is larger than one of its neighbors. Since it is smaller than `a[0]`, it must be larger than `a[2]`. In this way, we obtain a decreasing sequence, and conclude that the array is sorted in reverse order. But then `a[n-1]` is less than its left neighbor, and it would have to be the local minimum. This is a contradiction. \square

I implemented an algorithm that is similar to binary search, computing the index of the middle element each time, and comparing it against its neighbors (2 comparisons). Since the intervals cut the search space in half each time, the number of compares used is $\sim 2 \lg n$.

The algorithm begins by checking whether the first or the last element is a local minimum. These are a special case because they are the only elements with a single neighbor. If the array has only two elements, then one of these must be the local minimum because all elements are distinct.

Suppose that the array had more than two elements. The algorithm proceeds by setting `lo` to 1 and `hi` to `a.length - 2`. At this point, `a[lo-1]` is larger than `a[lo]` because `a[lo-1]` is `a[0]`, which is not the local minimum, as determined by the previous check. Similarly, `a[hi+1]` is larger than `a[hi]` by the same reasoning. The algorithm continues its search for a local minimum in the index interval `lo..hi` by checking whether the element at index `mid = lo + (hi - lo) / 2` is a local minimum. There are three cases:

- (i) If `a[mid]` is larger than `a[mid+1]`, this means `a[mid]` is not a local minimum, but `a[mid+1]` *may* be a local minimum. We set `lo = mid + 1` and continue the search in `mid+1..hi`.
- (ii) Otherwise if `a[mid]` is larger than `a[mid-1]`, this means again that `a[mid]` is not a local minimum but `a[mid-1]` *may* be a local minimum, so we set `hi = mid - 1` and continue the search in the index interval `lo..mid-1`.
- (iii) Otherwise, we have found a local minimum.

The search loops this way and ends when $lo > hi$. To see that the algorithm is able to find a local minimum this way, note that it maintains the invariant that $a[lo-1] > a[lo] \ \&\& \ a[hi+1] > a[hi]$ is **true** at all stages. Indeed, the condition holds before entering the loop, and whenever cases (i) and (ii) update lo or hi , the invariant still holds. Suppose that the loop ends because $lo > hi$. This happens when $lo = hi + 1$. The invariant assures us that $a[lo-1] > a[lo] \ \&\& \ a[hi+1] > a[hi]$. Since $a[lo-1]$ is $a[hi]$, and $a[hi+1]$ is $a[lo]$, this would say that $a[hi] > a[lo] \ \&\& \ a[lo] > a[hi]$ is **true**, which is impossible because there is a total ordering for integers. We conclude that the condition $lo > hi$ can never be reached, and hence, the program will end by returning the index `mid` of a local minimum.

Exercise 19. *Local minimum of a matrix.* Given an n -by- n array `a[]` of n^2 distinct integers, design an algorithm that finds a *strict local minimum*: an entry `a[i][j]` that is strictly less than its neighbors. Internal entries have 4 neighbors; entries on an edge have 3 neighbors; entries on a corner have 2 neighbors. The running time of your program should be proportional to n in the worst case, which means you cannot afford to examine all n^2 entries.

Solution. I failed to obtain a linear solution. I implemented a linearithmic solution by using the hint on the book website. It is similar to my solution in Exercise 19. It works by finding the minimum entry in the middle row, and then checking against its neighbors in the same column. The row search space is cut in half each time, and the search ends when its size is 0. A similar mathematical analysis to 1.4.18 would show this algorithm works.

Exercise 20. *Bitonic search.* An array is *bitonic* if it is comprised of an increasing sequence of integers followed immediately by a decreasing sequence of integers. Write a program that, given a bitonic array of n distinct `int` values, determines whether a given integer is in the array. Your program should use $\sim 3 \lg n$ compares in the worst case. *Extra credit:* use only $\sim 2 \lg n$ compares in the worst case.

Solution. Suppose that `a[]` is bitonic and let `n` be its length. If the search key is `a[0]` or `a[n-1]`, then there's nothing else to do.

Otherwise, we search for the key in the interval $1..n-2$. Since all of the entries are distinct, there must be a maximum entry. Finding the maximum itself can also be done via binary search. We simply start with the interval $1..n-2$. If the middle entry is less than its predecessor, then we've gone too far and fallen into the decreasing sequence, so we continue our search on the left. Otherwise, if the key is less than its successor in the sequence, then we are in the increasing sequence, so we search on the right.

Suppose now that `i` is the index of the maximum entry. If the search key equals the maximum entry, then our search is done. Otherwise, note that the maximum satisfies $a[i] > a[i+1]$ and $a[i] > a[i-1]$. Since the array is bitonic, this means that entries in the range $0..i-1$ form the increasing sequence, and entries in the range $i+1..n-1$ form an decreasing sequence. Thus, we can continue the search for the key on the disjoint sequences that do not include index `i` of the maximum by using binary search.

Exercise 1.4.21. *Binary search on distinct values.* Develop an implementation of binary search for `StaticSETofInts` (see page 99) where the running time of `contains()` is guaranteed to be $\sim \log d$, where d is the number of different integers in the array given as argument to the constructor.

Exercise 1.4.23. *Binary search with duplicates.* Modify `indexOf()` (page 47) so that it returns the index of the *first* (or *last*) entry equal to the query key. The order of growth of your method should be $\log n$ in the worst case, even if there are many entries in the array equal to the query key.

Solution. This is effectively the same as the same as Exercise 1.4.10, so I will omit the code.

Exercise 1.4.26. *3-collinearity.* Suppose that you have an algorithm that takes as input n distinct points in the plane and can return the number of triples that fall on the same line. Show that you can use this to solve the 3-sum problem. *Strong hint:* Use algebra to show that (a, a^3) , (b, b^3) , and (c, c^3) are collinear if and only if $a + b + c = 0$.

Solution.

Proof. We can begin by proving the claim in the hint. Suppose that a, b, c are distinct and (a, a^3) , (b, b^3) , and (c, c^3) are collinear. The line cannot be vertical because otherwise all x coordinates are equal, implying that a, b , and c are distinct, contrary to assumption.

Suppose that the points are on a line with slope m . Then

$$\begin{aligned} m &= \frac{c^3 - b^3}{c - b} = (c^2 + cb + b^2) \\ m &= \frac{c^3 - a^3}{c - a} = (c^2 + ca + a^2) \\ m &= \frac{b^3 - a^3}{b - a} = (b^2 + ba + a^2) \end{aligned}$$

where we have divided by common factors because all three integers are distinct. Equating the first two equations gives

$$\begin{aligned} c^2 + cb + b^2 &= c^2 + ca + a^2 \\ cb + b^2 &= ca + a^2 \\ b^2 - a^2 &= ca - cb \\ (b - a)(b + a) &= c(a - b) \\ b + a &= -c \\ a + b + c &= 0 \end{aligned}$$

Again we were able to divide by the common factor of $b - a$ because the integers are all distinct. We have shown that $a + b + c = 0$. By working backwards (retracting steps) we could show the converse.

Now, given this algorithm and the previously proven result, we could solve the 3-sum problem by creating, n triples of the form (α, α^3) , where α is one of the distinct n numbers. When the algorithm finds the there collinear points, we know their x -coordinates sum to 0. \square

Exercise 1.4.37. *Autoboxing performance penalty.* Run experiments to determine the performance penalty on your machine for using autoboxing and unboxing. Develop an implementation `FixedCapacityStackOfInts` and use a client such as `DoublingRatio` to

compare its performance with the generic `FixedCapacityStack<Integer>`, for a large number of `push()` and `pop()` operations.

Solution. For large n , the version that requires autoboxing to `Integer` requires around 50 times as long to run as the primitive `int` version.

Exercise 1.4.43. *Resizing array versus linked lists.* Run experiments to validate the hypothesis that resizing arrays are faster than linked lists for stacks (see Exercise 1.4.35 and Exercise 1.4.36). Do so by developing a version of `DoublingRatio` that computes the ratio of the running times of the two programs.

Solution. I observed that resizing implementation tends to run about 3 times as fast as the linked list implementation.

References

- [SW11] Robert Sedgewick and Kevin Wayne. *Algorithms*. 4th ed. Addison-Wesley, 2011.
ISBN: 9780321573513.